

# Banding in Games: A Noisy Rant

(revision 5)


Mikkel Gjøl, Playdead  
 [@pixelmager](https://twitter.com/pixelmager)

online link to this presentation  
[http://loopit.dk/banding\\_in\\_games.pdf](http://loopit.dk/banding_in_games.pdf)

see also  
<http://loopit.dk/publications/>

online link to this presentation [http://loopit.dk/banding\\_in\\_games.pdf](http://loopit.dk/banding_in_games.pdf)


- not going to go into too much detail about the theoretical background, try to give intuitive understanding and focus on issues



First a couple of example to motivate that this is actually still an issue in games.

amazing game, pixeljunk eden (if you haven't played it, you should literally leave the room and do so now!)

<http://pixeljunk.jp/library/Eden/>



so what we're looking at here are the “abrupt changes between shades of the same colour”



Skyrim, amazing do-whatever-you-want game... which means most will remember  
skyrim looking something like this

<http://www.elderscrolls.com/skyrim>



CONTINUE ↗

NEW

LOAD

CREDITS

QUIT



I was working on banding at the time I played it, so I remember the menus



...so I remember it looking mostly like this :)



...this is another beautiful indie-game, kentucky route zero


<http://kentuckyroutezero.com/>



which on most high-quality monitors will have quite noticeable banding (but almost get away with it due to their graphical style)

# What happens to your precious pixels?

(monitor expects gamma-corrected srgb)



the roughest pipeline-overview you will ever see  
things to note:

- monitor expects an input that is “roughly” srgb - in order to give the monitor srgb without terrible loss of precision, we either need >8bit precision, or need for all other stages to be aware that color-values end up in sRGB space... there are many interesting ways in which this can fail.

note also that blending reads from the framebuffer

see e.g. [http://www.opengl.org/registry/specs/ARB/framebuffer\\_sRGB.txt](http://www.opengl.org/registry/specs/ARB/framebuffer_sRGB.txt)

## Gamma-incorrectness exposes banding

- so just render gamma-correct...

Might not be so easy...

- sRGB may not be available and functional
- ...under all circumstances, RTT, blending etc
- ...on all your target platforms
- ...including mobile and old PC-hardware

(doable on most platforms, but there may be good reasons why your particular engine is not fully gamma-correct)

When incorrect the monitor will still correct the signal as if the image was in gamma-space, exposes all the horrible bands in the dark areas

If you use sRGB correctly, you're doing pretty well - you will generally hardly notice banding (though dark areas remain)

If you are not on a platform where it's readily available, or you want to get rid of the last issues, the rest of this presentation is for you

rtt => render to texture



KEEP  
CALM  
IT'S NOT A TALK  
ABOUT  
SRGB

...just wanted to go over this, so we know where we are in the pipeline, and there are good reasons a multiplatform engine could potentially have issues... mostly a last-gen / mobile issue though.

( thanks <http://www.keepcalm-o-matic.co.uk/> )



# KEEP CALM

...go read Steve Smith's excellent talk  
“Picture Perfect: Gamma through the Rendering Pipeline” from Gamefest 2007

<http://download.microsoft.com/download/b/5/b55d67ff-f1cb-4174-836a-bbf8f84fb7e1/Picture%20Perfect%20-%20Gamma%20Through%20the%20Rendering%20Pipeline.zip>


- someone else already did that way better:

<http://download.microsoft.com/download/b/5/b55d67ff-f1cb-4174-836a-bbf8f84fb7e1/Picture%20Perfect%20-%20Gamma%20Through%20the%20Rendering%20Pipeline.zip>

(mirror: [http://loopit.dk/Gamma\\_Through\\_the\\_Rendering\\_Pipeline.zip](http://loopit.dk/Gamma_Through_the_Rendering_Pipeline.zip) )

also <http://www.poynton.com/GammaFAQ.html>

“Don’t you just dither?”



This is the happy face of Thomas Rued, a man of experience who knows his pixels!  
...when all else fails, add noise

# Areas affected by color-banding

(...anything that renders gradients on screen)

Lighting

Fog

Alpha Blending

- additive (multipass rendering)
- subtractive
- multiplicative (lerp)
- ...creative :p

Particles

Posteffects, e.g. glow, aa, screen fades


(we will not go into undersampling and compression issues)

this talk is about the specific issues related to banding, where they occur and how to remedy them

..most noticeable on anything creating gradients: fog, falloffs, glow, particle splotches, from lights etc.

Noisy 'signals' (e.g. with a texture) "hides" most banding.

# Dithering: The art of noise



a variety of ways to do dithering, ...they all have to do with how the eye blurs out small changes (or integrates over an area). Squint a bit at the above image and you can barely tell the difference...

<https://www.shadertoy.com/view/MsIGR8>

...you should of course dither r,g,b separately for improved luminance

Dither Spatially (take advantage of eye accumulating multiple pixels)

Dither Temporally (take advantage of eye accumulating over time) => change dither pattern per frame

Scanline dither: x2 steps

2x2 ordered: x4 steps

8x8 ordered: x8(?) steps

# Ordered dithering on pixeldrugs (gamedevs... sheesh!)




...or other weirdness if you want to be ‘creative’ ;)  
(pixelartists)

- this is taken from a tutorial in drawing pixel-graphics with photoshop  
<http://danfessler.com/blog/hd-index-painting-in-photoshop>

## Much better dithering schemes exist


- Floyd Steinberg and related methods
- Most are error-diffusion based, don't map well to GPUs
- uniform noise is good enough for Pixar  
(and we don't mind a bit of grain in our camera to get less sterile images)




error diffusion based methods rely on knowledge of the surrounding area, something that doesn't map well to the GPU, which prefer doing a lot of unrelated things in parallel

<http://www.realtimerendering.com/blog/2011-color-and-imaging-conference-part-ii-courses-a/>

<http://www.joesfer.com/?p=149>  
<http://www.joesfer.com/?p=108>




unity lightbuffer - already in a logarithmic colorspace to enable HDR lighting, so no sense in additionally using srgb



modified Internal-PrePassLighting.shader to do (monochrome) random dithering  
...in order to not create discontinuities at deferred light-edge (light fades out to edge),  
we subtract from the intensity instead of adding...  
..so slightly darker, but otherwise uniform (we vary noise over time as well, so with a  
bad monitor, and/or eyesight, the result is quite good)



# quantizing signal by truncation



doing simple quantization (by truncating) is always darker than intended signal (by 0.5bit on avg)

# quantizing signal, error

accumulated err 0.5LSB, max err 1LSB, 0.5LSB bias towards darker



signal  
 $\text{floor}(\text{signal})$



error rounding down

max error of 1 a Least-Significant-Bit

0.5 LSB error on avg

0.5 LSB bias towards darker result


# rounding



signal  
signal+0.5  
 $\text{floor}(\text{signal}+0.5)$

# rounding

accumulated err is 0, max err 0.5LSB, no bias





Though the average error is 0.25LSB, the accumulated error is 0 (intuitively as area rounding up == area rounding down), but absolute error at any value is as much as 0.5\*LSB

Signal does not lose 'energi' as opposed to trunc  
Smaller err, smaller maxerr (half before),

now at least on average, it's the same signal as the original. It is unbiased.

# rounding+signed random

(white noise, uniformly distributed [-0.5;0.5[ LSB )



signal  
signal+0.5

(signal+0.5) + [-0.5;0.5[  
= signal + [0.0;1.0[


adding a uniform signed noise means we're still not biasing the signal - sometimes brighter, sometimes darker, so on average it does neither...

...property still valid that avg error is 0

This is exactly the same as adding a normalized [0;1[ to the signal

## “area” of noise


accumulated error still 0, max error 1LSB, no bias



max err of course now larger as we just added noise (but still: no bias)  
intuitively: on average quantizing as much to value above as value below, so end up with the same signal... on average

- same as before, on average the positive/negative noise will cancel out, and we get the original signal on average.  
Error still varies (intuitively: When signal crosses the truncated value, the resulting error is 0)  
...we get slightly worse max error than from just rounding.. but....

# magnitude of uniform noise



...if we look at the accumulated error for a single value (integrating over the red/blue stippled line), the error will now cancel out as well as for entire signal, resulting in the original signal for any arbitrary single value

This is a property of the noise being uniformly distributed

intuitively, since the noise-distribution is uniform, when integrating across the line shown the length of the line corresponds to the probability that the value will either round up or down



...integrating  $\text{floor}(f)$  /  $\text{ceil}(f)$  across this line, we'll end up at the signal

**...noise is uniformly distributed, but the resulting visual noise is NOT uniformly distributed - almost no noise near “correct” values**

**Really no reason to add noise where signal -is- its truncated value**

# triangular noise distribution

$\text{floor}(\text{signal} + \text{hash(seed1)} + \text{hash(seed2)} - 0.5) \Rightarrow [-0.5; 1.5[$



HINDSIGHT: GPUs round, so  
dither-range should be  $[-1; 1]$   
i.e.  $\text{hash}(s1) + \text{hash}(s2) - 1.0$



Supposedly a PDF with better characteristics, commonly used for audio-dithering. Effectively adds noise in low-noise-areas, giving a more uniform noise-appearance.

We're not currently using this, wikipedia suggests using it if the result is to be "worked on further"...

Visually uniform noise in gradient (pretty much just adds noise to low-error areas)

HINDSIGHT: The phenomenon is called "noise modulation". This turns out to be very important to achieve a banding-free look, and we relied heavily on it in the shipped product.

HINDSIGHT: instead of calculating two hashes from scratch, they can be combined to make them a lot faster. Also, remapping a single hash to have a triangular PDF may be faster, see <https://www.shadertoy.com/view/4t2SDh>

## references


<http://en.wikipedia.org/wiki/Dither>

[http://www.ece.rochester.edu/courses/ECE472/resources/Papers/Lipshitz\\_1992.pdf](http://www.ece.rochester.edu/courses/ECE472/resources/Papers/Lipshitz_1992.pdf)

<https://www.shadertoy.com/view/4ssXRX>

# various distributions

HINDSIGHT:  
blue noise is less noticeable



“Awesome, if triangular is good, gaussian must be even betterer!” (spoiler: don’t do it)

From left to right

- A uniform distribution in the interval  $[0;1[$  creates areas of little to no noise, giving the impression of a smooth signal, but non-uniform noise.
- expanding the noise to  $[-0.5;1.5[$  creates areas of too much noise (in the “overlapping” regions)
- Switching to a triangular distribution  $[-0.5;1.5[$  gives a nice uniform distribution (the “overlapping” noise-regions accumulate to 1)
- Using a gaussian distribution again creates areas with too little noise
- using a gaussian-pdf in range  $[-1.0;2.0[$  appears very similar to triangular, but with more noticeable noise

Empiric conclusion: Use a triangular distribution in the interval  $[-0.5;1.5[$

HINDSIGHT: The reason a triangular distribution is “enough”, is that it does not exhibit “noise modulation”. Neither do gaussians, but this is the main benefit.

HINDSIGHT: GPUs round, so dither-range should be  $[-1;1[$

HINDSIGHT: Using high-pass filtered noise can give even better results, as the noise is less noticeable. This is what is called “noise shaping” in audio-dithering.

# various distributions, errors





uniform  
[0.0;1.0[

uniform  
[-0.5;1.5[


triangular  
[-0.5;1.5[

gaussian  
[-0.5;1.5[

gaussian  
[-1.0;2.0[



HINDSIGHT: this shows the absolute value of the error,  $\text{abs}(\text{err})$ , which makes it hard to see that positive/negative errors “cancel each out” in the noisy areas. Remapping the actual error to a [0;1] range gives a much more accurate depiction of the error.





Example of dithering using triangular noise [-0.5;1.0[ vs. uniform noise [0;1] - triangular noise makes the noise more evident, but also more uniform - no areas with little noise.

...the effect is certainly subtle, but better is better :) (and also more noisy :p )


# subtracting noise

same properties, with an offset



just a quick note that subtracting noise of course results in a signal with the same properties, but offset

## Dithering after quantisation does not remove banding (e.g. LDR film-grain at end of frame)



"lets just add grain at end of frame to make banding go away"

...only works if input signal is kept at high precision, so e.g. HDR-rendering and adding grain at time of tonemapping is fine...

Otherwise you are just adding noise on top of an already banding image.

# the right color space

```
vec4 fragmain(vec2 fragpos) {  
    vec4 outcol;  
    ...  
    return outcol;  
}
```

...spend a bit on time on this, because it's important  
colors are in linear-space in the pixelshader (or should be unless you're messing  
something up)...

In essense: Dither just before quantizing.

# the right color space (linear target, dithered)

```
vec4 hash42n( vec2 seed ) { .. }

vec4 ditherRGBA( vec4 c, vec2 seed ) {
    return c + hash42n( seed ) / 255.0; //8bit
}
vec4 fragmain(vec2 fragpos) {
    vec4 outcol;
    ...
    return ditherRGBA( outcol, pos );
}
```

...hash42n is a function returning a vec4 of uniformly distributed [0;1[ random numbers given a vec2 as seed.

It can be ALU or a texture-lookup depending on where you have most to spare...

# the right color space (manual srgb rendertarget)

```
vec4 ditherRGBA( vec4 c, vec2 seed ) {
    return c + hash42n( seed ) / 255.0; //8bit
}
vec4 fragmain(vec2 fragpos) {
    vec4 outcol;
    ...
    return ditherRGBA(lin2srgb(outcol), fragpos);
}
```

convert first, dither afterwards... dither in whatever space the color will be quantized in.

# the right color space

## (manual logarithmic buffer, e.g. unity lightbuffer)

```
vec4 ditherRGBA( vec4 c, vec2 seed ) {  
    return c + hash42n( seed ) / 255.0; //8bit  
}  
vec4 fragmain(vec2 fragpos) {  
    vec4 outcol;  
    ...  
    return ditherRGBA(exp2(-outcol), fragpos);  
}
```

again, do it as close to quantization as possible - jittering the light-value does not help...

# the right color space (autoconverting srgb rendertarget)

```
vec4 ditherRGBA(vec4 c, vec2 s) { return c+hash42n(s)/255.0; }
vec4 fragmain(vec2 fragpos) {
    vec4 outcol;
    ...
    return srgb2lin(ditherRGBA(lin2srgb(outcol), fragpos));
}
```

...need to dither in whatever space the color will be quantized in.


# the right color space (autoconverting srgb rendertarget)

```
vec4 ditherRGBA(vec4 c, vec2 s) { return c+hash42n(s)/255.0; }
vec4 fragmain(vec2 fragpos) {
    vec4 outcol;
    ...
    return srgb2lin(ditherRGBA(lin2srgb(outcol), fragpos));
}
//doesn't strictly need to be all that accurate for
this...
vec4 srgb2lin(vec4 c) { return vec4(c.rgb*c.rgb, c.a); }
vec4 lin2srgb(vec4 c) { return vec4(sqrt(c.rgb), c.a); }
```


$$(\sqrt{c}+n)^2 == c + 2*n*\sqrt{c} + n^2$$

...but pow( sqrt( c ) + rnd/255.0, 2.0 ); is probably faster :p

... see also Timothy Lottes': <https://www.shadertoy.com/view/Md2XWw>



<https://www.shadertoy.com/view/XsfXzf>




menu of steam big picture (best thing that has happened to pc-gaming for a long time)  
- has noise added, but interestingly in what appears to be the wrong color-space.

# Triangular dithering at boundaries

HINDSIGHT

It turns out, that due to clamping, the boundaries around black/white (0/1) become wrong when using 2LSB triangular dithering.




<https://computergraphics.stackexchange.com/questions/5904/whats-a-proper-way-to-clamp-dither-noise/5952#5952>

intuition: As the noise gets clamped, clamping destroys the symmetry in the noise - and for that reason the average does not tend towards the actual signal

# Triangular dithering at boundaries

HINDSIGHT

A solution is to lerp to a 1LSB uniform-PDF noise at boundaries



<https://computergraphics.stackexchange.com/questions/5904/whats-a-proper-way-to-clamp-dither-noise/5952#5952>

# Triangular dithering at boundaries

HINDSIGHT

```
vec2 rnd = hash22(seed);

float dithertri = (rnd.x + rnd.y - 1.0); //note: symmetric, triangular dither, [-1;1[

float dithernorm = rnd.x - 0.5; //note: symmetric, uniform dither [-0.5;0.5[

float sizt_lo = clamp( v/(0.5/7.0), 0.0, 1.0 );

float sizt_hi = 1.0 - clamp( (v-6.5/7.0)/(1.0-6.5/7.0), 0.0, 1.0 );

dither = lerp( dithernorm, dithertri, min(sizt_lo, sizt_hi) );
```

<https://computergraphics.stackexchange.com/questions/5904/whats-a-proper-way-to-clamp-dither-noise/5952#5952>

# alpha blending

## blend-modes

- additive,  $c = c0 + c1; //\text{dither } c0, c1$
- subtractive,  $c = c0 - c1; //\text{dither } c0, c1$
- multiplicative,  $c = c0 * c1; //... :()$
- creative :p

adding / subtracting does not change the magnitude of LSB, so the dithering is fine.  
You still end up adding 1LSB of noise for each input, which adds up...  
multiplying ruins the whole thing - no way to determine the magnitude of a LSB is  
after multiplication of “some value [0;1]” with a constant (accessing the  
destination-buffer of course is even worse)

# alphablending, non-constructive rnd (blending multiple shaders to same pixel)

```
vec4 ditherRGBA( vec4 c, vec2 seed ) {  
    return c + hash42n( seed ) / 255.0; //8bit  
}  
  
vec4 fragmain(vec2 fragpos) {  
    vec4 outcol;  
    ...  
    vec2 UNIQUE_SEED = vec2(0.6849);  
    return ditherRGBA( outcol, fragpos+UNIQUE_SEED );  
}
```

...blending tends to deal with translucency, so

...if blending multiple objects with to the same pixel, using a different noise-value for each layer helps reduce noise.

# alphablending, non-constructive rnd (blending same shader to same pixel, particle systems)

```
vec4 ditherRGBA( vec4 c, vec2 seed ) {  
    return c + hash42n( seed ) / 255.0; //8bit  
}  
  
vec4 fragmain(vec2 fragpos) {  
    vec4 outcol;  
    ...  
    vec2 UNIQUE_SEED = vec2(0.6849 + vspos.z);  
    return ditherRGBA( outcol, fragpos+UNIQUE_SEED );  
}
```

...if blending multiple objects with the same **shader**, using a different noise-value for each layer helps reduce noise.

E.g. a particle-system with many layers could use the viewspace-z as a seed on top of screenposition.



## further blending issues

Anything that multiplies something onto **dstcol**

- do as much in shader as possible
  - don't use blendunit for srcAlpha if you can do it
- 2pass: subtract from dst, then add color  
(doesn't band... but also doesn't look the same)
- add **moar noise** to dst-alpha :(  
...helps mask problem, but not a "solution".

# multiplicative blending

(dithered dstcolor) \* srcalpha, can't dither post-blend



0.5 LSB more noise required  
(fine if added to incoming signal at highp)

scales down noise, need to add noise again before quantization

...really, we would be fine if the noise added to src-color would just remain at high precision until after blending (more on that in a second...)

# alphablending (prefer premultiplied alpha)

```
vec4 ditherRGBA( vec4 c, vec2 seed ) {  
    return c + hash42n( seed ) / 255.0; //8bit  
}  
  
vec4 fragmain(vec2 fragpos) {  
    vec4 outcol;  
    ...  
    outcol.rgb *= outcol.a;  
    return ditherRGBA( outcol, fragpos );  
}
```

additive blending with SRC\_ALPHA means adding noise before multiplying!

premult: then at least the incoming signal is ok...


dithering alpha depends on how it's used in the blend-function... prefer  
pre-multiplied alpha to SRCALPHA, ONE\_MINUS\_SRCALPHA

[http://home.comcast.net/~tom\\_forsyth/blog.wiki.html#\[\[Premultiplied%20alpha\]\]](http://home.comcast.net/~tom_forsyth/blog.wiki.html#[[Premultiplied%20alpha]])


]

...allows to blend colors smaller than 1/255 (as they are dithered afterwards)

alpha-noise:  
1xLSB



alpha-noise:  
17xLSB




...empirical values \o/ Set manually by artists

# precision of blend-unit?

it's complicated... and platform specific

HINDSIGHT:  
consistent with d3d11-spec



GPU vendor1  
lowp (or quantize before blend?)      GPU vendor2  
highp (or quantize after blend?)

honestly not entirely sure what is going on here... but seems vendor2 runs blending at higher precision (or quantizes after blending). Fortunately, vendor2 can be found in all major new consoles.

$a * \text{srccol} + b * \text{dstcol}$

$\text{trunc}(a) * \text{trunc}(c0) + \text{trunc}(b) * c1t // \dots :$   
 $\text{trunc}(c0d + b * c1t) //$  dithering-noise gets added "correctly"

-  
**Hindsight:** This appears to be consistent with the d3d11 spec, which states that blending only has to be performed at rendertarget-precision (e.g.  $\text{rgba8} \Rightarrow 8\text{bit}$  blending precision):

*"When a RenderTarget has a fixed point format <snip> blending operations may be performed at equal or more (e.g up to float32) precision/range than the output format."*

# ~~programmable blending~~

- nvidia opengl extension (also rsx on ps3)  
GL\_NV\_texture\_barrier
- intel dx11-extension pixel shader ordering
- mobile tiled, defered architectures,  
opengles  
GL\_EXT\_shader\_framebuffer\_fetch
- Currently no unified solution (dx12?)

way overkill - on platforms that support this, there are likely better alternatives: sRGB, higher precision buffers etc.

nvidia / ps3 texture barrier [http://www.opengl.org/registry/specs/NV/texture\\_barrier.txt](http://www.opengl.org/registry/specs/NV/texture_barrier.txt)  
Intel pixelsync

<https://software.intel.com/en-us/blogs/2013/03/27/programmable-blend-with-pixel-shader-ordering>

moar pixelsync

<http://advances.realtimerendering.com/s2013/2013-07-23-SIGGRAPH-PixelSync.pdf>


Nvidia Raster-ordered view

[http://international.download.nvidia.com/geforce-com/international/pdfs/GeForce\\_GT\\_X\\_980\\_Whitepaper\\_FINAL.PDF](http://international.download.nvidia.com/geforce-com/international/pdfs/GeForce_GT_X_980_Whitepaper_FINAL.PDF)

# post-processing

- Glow, screen-fades, vignettes...
- Typically happens on 8bit textures  
(for very good bandwidth reasons)
- Writes images in multiple steps
  - ...so you should dither each step!  
(or just the “important” ones, ymmv)
  - if need to match former incorrect, truncated,  
non-dithered color, add (signed)  $\text{rand}[-0.5;0.5]$   
instead of (normalized)  $\text{rand}[0;1]$

## post-processing: Glow



our solution dithers after every pass

manually converts all intermediate rendertargets to srgb (pow2, nothing fancy)

results in larger than 1pixel noise, but still not visible


sRGB goes a long way! adding noise was necessary for us though




limbo was entirely grayscale, also through the rendering pipeline  
rendered in 10:10:10:2 on pc/xbox360, rgba16 on ps3 - color “offset-dithering” on  
tonemapping (plus tons of film-grain in post!)  
- please keep this “entirely gray-scale” image in mind, for when you during the next  
slide think “oh this can’t possibly work!”



color-offset dithering  
idea by Kim Steen Riber, now at Unity





Very fast. Works for grayscale ONLY

“blueish” - color offset used in Limbo, adds x3 luminance steps

“red/green-ish” - weighted by rgb individual luminances, still only x3 steps  
(even better using a custom LUT for luminances:

<https://www.shadertoy.com/view/4sjXWw> , adds ~x16 steps)

dithering overview: <https://www.shadertoy.com/view/MsIGR8>

# concluding recommendations


- Use highest affordable precision available (duh)
- Use sRGB if feasible across target platforms
- Dither using triangular noise [-0.5;1.0[ during quantisation  
(e.g. at tonemapping, or **any** write to an 8bit rt)  
HINDSIGHT: GPUs round, so dither-range should be [-1;1[
- Add noise in the right color-space  
(the color-space your quantised values are stored in)

# welcome side-effects from dithering

- Can now use and blend colors < 1/255
- No color-shifts in gradients (rgb truncate differently)
- Tons of additive layers? Now adds noise instead of banding
- More colors perceived than the monitor is capable of
- smoooooth screen-fades!

more or less intended side-effects

ooh, this stuff works for gbuffer-normals too!



<https://www.youtube.com/watch?v=h59LwyJbfzs>

**“I wonder what that sounds like?”**

(said no-one ever... and the sound-guy at Playdead)

<https://www.youtube.com/watch?v=h59LwyJbfzs>

not something you hear about graphics algorithms every day :)

# Thanks

Daniel Povlsen, [@danielpovlsen](#)

Mikkel Svendsen, [@lkarosAV](#)

Jakob Schmid, [@jakobschmid](#)

All images are copyright of their respective authors

# references

<http://en.wikipedia.org/wiki/Dither>

<http://caca.zoy.org/wiki/libcaca/study/introduction>

<http://sandervanrossen.blogspot.dk/2012/02/hdr-dithering.html>

<http://www.realtimerendering.com/blog/2011-color-and-imaging-conference-part-ii-courses-a/>

<https://bartwronski.com/2016/10/30/dithering-part-one-simple-quantization/>

<https://www.shadertoy.com/view/MsIGR8> - various dithering methods

<https://www.shadertoy.com/view/4dfXW8> - example-workspace used in this presentation

<http://xiph.org/video/vid2.shtml> - on audio dithering

[http://www.ece.rochester.edu/courses/ECE472/Site/Assignments/Entries/2009/1/15\\_Week\\_1\\_files/Lipshitz\\_1992.pdf](http://www.ece.rochester.edu/courses/ECE472/Site/Assignments/Entries/2009/1/15_Week_1_files/Lipshitz_1992.pdf) - dithering theory

<http://download.microsoft.com/download/b/5/b55d67ff-f1cb-4174-836a-bbf8f84fb7e1/Picture%20Perfect%20-%20Gammma%20Through%20the%20Rendering%20Pipeline.zip>

# PLAYDEAD IS HIRING



[job@playdead.com](mailto:job@playdead.com)

~~bad~~ bonus slides!

## **ARB\_framebuffer\_sRGB (2008)**

“...render into a framebuffer that is scanned to a monitor configured to assume framebuffer color values are sRGB encoded. This assumption is roughly true of most PC monitors with default gamma correction. This allows applications to achieve faithful color reproduction for OpenGL rendering without adjusting the monitor's gamma correction.”

[http://www.opengl.org/registry/specs/ARB/framebuffer\\_sRGB.txt](http://www.opengl.org/registry/specs/ARB/framebuffer_sRGB.txt)


# HDR resolve




<http://sandervanrossen.blogspot.dk/2012/02/hdr-dithering.html>

# Real World Assets

(it's just noise, not magic)




<http://www.sandraandwoo.com/comics/2012-11-19-0430-software-engineering-now-with-cats.png>



bugs from the real world

- a lightsource brighter than the sun, illuminating dark, black granite stone.



bugs from the real world (real-world assets)

"what the... oh"....not banding, just faceted geometry...